

# LeetCode SQL Problem Solving Questions With Solutions

## LeetCode SQL Solutions

### 175. Combine Two Tables | Easy | [LeetCode](#)

Table: Person

TEXT

Column Name	Type
PersonId	int
FirstName	varchar
LastName	varchar

PersonId is the primary key column for this table.

### Table: Address

```
TEXT
+-----+-----+
| Column Name | Type   |
+-----+-----+
| AddressId   | int    |
| PersonId    | int    |
| City         | varchar |
| State        | varchar |
+-----+-----+
```

AddressId is the primary key column for this table.

Write a SQL query for a report that provides the following information for each person in the Person table, regardless if there is an address for each of those people:

```
TEXT
FirstName, LastName, City, State
```

## Solution

```
SQL
SELECT p.FirstName, p.LastName, a.City, a.State
FROM Person p
LEFT JOIN Address a
ON p.PersonId = a.PersonId;
```

## 176. Second Highest Salary | Easy | [LeetCode](#)

Write a SQL query to get the second highest salary from the Employee table.

```
TEXT
+----+-----+
| Id | Salary |
+----+-----+
| 1  | 100   |
| 2  | 200   |
| 3  | 300   |
+----+-----+
```

For example, given the above Employee table, the query should return 200 as the second highest salary. If there is no second highest salary, then the query should return null.

```
TEXT
+-----+
| SecondHighestSalary |
+-----+
| 200                 |
+-----+
```

## Solution

```
SQL
#Solution 1:
SELECT Max(Salary) SecondHighestSalary
FROM Employee WHERE Salary < (SELECT MAX(Salary) FROM Employee)

#Solution 2:
WITH CTE AS (SELECT DISTINCT Salary
FROM Employee
ORDER BY Salary DESC
LIMIT 2)

SELECT Salary as SecondHighestSalary
FROM CTE
ORDER BY Salary Asc
LIMIT 1;
```

```
#Solution 3:

WITH CTE AS
(
    SELECT Salary,
           DENSE_RANK() OVER (ORDER BY Salary DESC) AS DENSERANK
      FROM Employee
)
SELECT Salary SecondHighestSalary
  FROM CTE
 WHERE DENSERANK = 2;
```

## 177. Nth Highest Salary | Medium | [LeetCode](#)

Write a SQL query to get the nth highest salary from the Employee table.

TEXT
+-----+   Id   Salary   +-----+   1    100       2    200       3    300     +-----+

For example, given the above Employee table, the nth highest salary where n = 2 is 200. If there is no nth highest salary, then the query should return null.

TEXT
+-----+   getNthHighestSalary(2)   +-----+   200                   +-----+

## Solution

SQL 

```
CREATE FUNCTION getNthHighestSalary(N INT) RETURNS INT
BEGIN
    SET N = N-1;
    RETURN(
        SELECT DISTINCT Salary FROM Employee ORDER BY Salary DESC
        LIMIT 1 OFFSET N
    );
END
```

## 178. Rank Scores | Medium | [LeetCode](#)

Write a SQL query to rank scores. If there is a tie between two scores, both should have the same ranking. Note that after a tie, the next ranking number should be the next consecutive integer value. In other words, there should be no “holes” between ranks.

TEXT

Id	Score
1	3.50
2	3.65
3	4.00
4	3.85
5	4.00
6	3.65

For example, given the above `scores` table, your query should generate the following report (order by highest score):

TEXT

score	Rank
4.00	1
4.00	1
3.85	2
3.65	3
3.65	3
3.50	4

**Important Note:** For MySQL solutions, to escape reserved words used as column names, you can use an apostrophe before and after the keyword. For example `Rank`.

## Solution

SQL

```
SELECT score, DENSE_RANK() OVER (ORDER By Score DESC) AS "Rank"  
FROM Scores;
```



## 180. Consecutive Numbers | Medium | [LeetCode](#)

Table: Logs

TEXT

Column Name	Type
id	int
num	varchar

`id` is the primary key for this table.

Write an SQL query to find all numbers that appear at least three times consecutively.

Return the result table in any order.

The query result format is in the following example:

TEXT

Logs table:

Id	Num
1	1
2	1
3	1
4	2
5	1
6	2
7	2

Result table:

ConsecutiveNums
1

1 is the only number that appears consecutively for at least three times.

## Solution

SQL

```
SELECT a.Num as ConsecutiveNums
FROM Logs a
JOIN Logs b
ON a.id = b.id+1 AND a.num = b.num
JOIN Logs c
ON a.id = c.id+2 AND a.num = c.num;
```



## 181. Employees Earning More Than Their Managers | Easy | LeetCode

The `Employee` table holds all employees including their managers. Every employee has an `Id`, and there is also a column for the manager `Id`.

TEXT

Id	Name	Salary	ManagerId
1	Joe	70000	3
2	Henry	80000	4
3	Sam	60000	NULL
4	Max	90000	NULL

Given the `Employee` table, write a SQL query that finds out employees who earn more than their managers. For the above table, Joe is the only employee who earns more than his manager.

TEXT

Employee
Joe

## Solution

SQL

```
SELECT E.Name as "Employee"  
FROM Employee E  
JOIN Employee M
```



6 | 0

```
ON E.ManagerId = M.Id  
AND E.Salary > M.Salary;
```

## 182. Duplicate Emails | Easy | [LeetCode](#)

Write a SQL query to find all duplicate emails in a table named `Person`.

```
TEXT  
+----+-----+  
| Id | Email |  
+----+-----+  
| 1  | a@b.com |  
| 2  | c@d.com |  
| 3  | a@b.com |  
+----+-----+
```

For example, your query should return the following for the above table:

```
TEXT  
+-----+  
| Email |  
+-----+  
| a@b.com |  
+-----+
```

**Note:** All emails are in lowercase.

### Solution

```
SQL  
#Solution- 1:  
SELECT Email  
FROM Person  
GROUP BY Email  
HAVING count(*) > 1
```



```

#Solution- 2:

WITH CTE AS(
    SELECT Email, ROW_NUMBER() OVER(PARTITION BY Email ORDER BY Email) AS RN
    FROM Person
)

SELECT Email
FROM CTE
WHERE RN > 1;

```

## 183. Customers Who Never Order | Easy | [LeetCode](#)

Suppose that a website contains two tables, the `Customers` table and the `Orders` table. Write a SQL query to find all customers who never order anything.

Table: `Customers`.

```

TEXT
+----+-----+
| Id | Name  |
+----+-----+
| 1  | Joe   |
| 2  | Henry |
| 3  | Sam   |
| 4  | Max   |
+----+-----+

```

Table: `orders`.

```

TEXT
+----+-----+
| Id | CustomerId |
+----+-----+
| 1  | 3          |
| 2  | 1          |
+----+-----+

```

Using the above tables as example, return the following:

```
TEXT
+-----+
| Customers |
+-----+
| Henry      |
| Max        |
+-----+
```

## Solution

```
SQL
#Solution- 1:
SELECT Name AS Customers
FROM Customers
LEFT JOIN Orders
ON Customers.Id = Orders.CustomerId
WHERE CustomerId IS NULL;

#Solution- 2:
SELECT Name as Customers
FROM Customers
WHERE Id NOT IN(
    SELECT CustomerId
    FROM Orders
)
```



## 184. Department Highest Salary | Medium | [LeetCode](#)

The Employee table holds all employees. Every employee has an Id, a salary, and there is also a column for the department Id.

TEXT

Id	Name	Salary	DepartmentId
1	Joe	70000	1
2	Jim	90000	1
3	Henry	80000	2
4	Sam	60000	2
5	Max	90000	1

The Department table holds all departments of the company.

TEXT

Id	Name
1	IT
2	Sales

Write a SQL query to find employees who have the highest salary in each of the departments. For the above tables, your SQL query should return the following rows (order of rows does not matter).

TEXT

Department	Employee	Salary
IT	Max	90000
IT	Jim	90000
Sales	Henry	80000

### ***Explanation:***

Max and Jim both have the highest salary in the IT department and Henry has the

highest salary in the Sales department.

## Solution

SQL

```
SELECT Department.Name AS Department, Employee.Name AS Employee, Salary
FROM Employee
JOIN Department
ON Employee.DepartmentId = Department.Id
WHERE (DepartmentId, Salary) IN(
    SELECT DepartmentId, MAX(Salary) AS Salary
    FROM Employee
    GROUP BY DepartmentId
);
```



## 185. Department Top Three Salaries | Hard | [LeetCode](#)

The `Employee` table holds all employees. Every employee has an `Id`, and there is also a column for the department `Id`.

TEXT

Id   Name`   Salary   DepartmentId
1   Joe   85000   1
2   Henry   80000   2
3   Sam   60000   2
4   Max   90000   1
5   Janet   69000   1
6   Randy   85000   1
7   Will   70000   1

The `Department` table holds all departments of the company.

TEXT

Id	Name
1	IT
2	Sales

Write a SQL query to find employees who earn the top three salaries in each of the department. For the above tables, your SQL query should return the following rows (order of rows does not matter).

TEXT

Department	Employee	Salary
IT	Max	90000
IT	Randy	85000
IT	Joe	85000
IT	Will	70000
Sales	Henry	80000
Sales	Sam	60000

### ***Explanation:***

In IT department, Max earns the highest salary, both Randy and Joe earn the second highest salary, and Will earns the third highest salary. There are only two employees in the Sales department, Henry earns the highest salary while Sam earns the second highest salary.

## **Solution**

SQL

```
WITH department_ranking AS (
SELECT Name AS Employee, Salary ,DepartmentId
,DENSE_RANK() OVER (PARTITION BY DepartmentId ORDER BY Salary DESC) AS rnk
FROM Employee
```

```
)
```

```
SELECT d.Name AS Department, r.Employee, r.Salary
FROM department_ranking AS r
JOIN Department AS d
ON r.DepartmentId = d.Id
WHERE r.rnk <= 3
ORDER BY d.Name ASC, r.Salary DESC;
```

## 196. Delete Duplicate Emails | Easy | [LeetCode](#)

Write a SQL query to delete all duplicate email entries in a table named `Person`, keeping only unique emails based on its smallest Id.

TEXT

Id	Email
1	john@example.com
2	bob@example.com
3	john@example.com

`Id` is the primary key column for this table. For example, after running your query, the above `Person` table should have the following rows:

TEXT

Id	Email
1	john@example.com
2	bob@example.com

### Note:

Your output is the whole `Person` table after executing your sql. Use `delete`

statement.

## Solution

SQL

```
DELETE p2
FROM Person p1
JOIN Person p2
ON p1.Email = p2.Email
AND p1.id < p2.id
```



## 197. Rising Temperature | Easy | [LeetCode](#)

Table: Weather

TEXT

Column Name	Type
id	int
recordDate	date
temperature	int

id is the primary key for this table.

This table contains information about the temperature in a certain day.

Write an SQL query to find all dates' id with higher temperature compared to its previous dates (yesterday).

Return the result table in any order.

The query result format is in the following example:

6 | 0

TEXT

Weather

id   recordDate   Temperature
1   2015-01-01   10
2   2015-01-02   25
3   2015-01-03   20
4   2015-01-04   30

Result table:

+----+
id
+----+
2
4
+----+

In 2015-01-02, temperature was higher than the previous day (10 -> 25).

In 2015-01-04, temperature was higher than the previous day (20 -> 30).

## Solution

SQL



```
#Solution- 1:  
SELECT t.Id  
FROM Weather AS t, Weather AS y  
WHERE DATEDIFF(t.RecordDate, y.RecordDate) = 1  
AND t.Temperature > y.Temperature;
```

```
#Solution- 2:
```

```
SELECT t.Id  
FROM Weather t  
JOIN Weather y  
ON DATEDIFF(t.recordDate, y.recordDate) = 1 AND  
t.temperature > y.temperature;
```

## 262. Trips and Users | Hard | [LeetCode](#)

Table: Trips

TEXT

Column Name	Type
Id	int
Client_Id	int
Driver_Id	int
City_Id	int
Status	enum
Request_at	date

Id is the primary key for this table.

The table holds all taxi trips. Each trip has a unique Id, while Client\_Id and Driver\_Id are foreign keys to the Users table. Status is an ENUM type of ('completed', 'cancelled\_by\_driver', 'cancelled\_by\_client').

---

Table: Users

TEXT

Column Name	Type
Users_Id	int
Banned	enum
Role	enum

Users\_Id is the primary key for this table.

The table holds all users. Each user has a unique Users\_Id, and Role is an ENUM type of ('client', 'driver'). Banned is a boolean type. Status is an ENUM type of ('Yes', 'No').

---

Write a SQL query to find the cancellation rate of requests with unbanned users (both client and driver must not be banned) each day between "2013-10-01" and "2013-10-03".

The cancellation rate is computed by dividing the number of canceled (by client or driver) requests with unbanned users by the total number of requests with unbanned users on that day.

Return the result table in any order. Round Cancellation Rate to two decimal points.

The query result format is in the following example:

TEXT

Trips table:

Id	Client_Id	Driver_Id	City_Id	Status	Request_at
1	1	10	1	completed	2013-10-01
2	2	11	1	cancelled_by_driver	2013-10-01
3	3	12	6	completed	2013-10-01
4	4	13	6	cancelled_by_client	2013-10-01
5	1	10	1	completed	2013-10-02
6	2	11	6	completed	2013-10-02
7	3	12	6	completed	2013-10-02
8	2	12	12	completed	2013-10-03
9	3	10	12	completed	2013-10-03
10	4	13	12	cancelled_by_driver	2013-10-03

Users table:

Users_Id	Banned	Role
1	No	client
2	Yes	client
3	No	client
4	No	client
10	No	driver
11	No	driver
12	No	driver
13	No	driver

Result table:

Day	Cancellation Rate
2013-10-01	0.33
2013-10-02	0.00
2013-10-03	0.50

On 2013-10-01:

- There were 4 requests in total, 2 of which were canceled.
- However, the request with Id=2 was made by a banned client (User\_Id=2), so it is ignored.
- Hence there are 3 unbanned requests in total, 1 of which was canceled.
- The Cancellation Rate is  $(1 / 3) = 0.33$

On 2013-10-02:

- There were 3 requests in total, 0 of which were canceled.
- The request with Id=6 was made by a banned client, so it is ignored.
- Hence there are 2 unbanned requests in total, 0 of which were canceled.
- The Cancellation Rate is  $(0 / 2) = 0.00$

On 2013-10-03:

- There were 3 requests in total, 1 of which was canceled.
- The request with Id=8 was made by a banned client, so it is ignored.
- Hence there are 2 unbanned requests in total, 1 of which were canceled.
- The Cancellation Rate is  $(1 / 2) = 0.50$

## Solution

SQL



```
SELECT Request_at AS Day,
ROUND(SUM(IF(Status<>"completed", 1, 0))/COUNT(Status),2) AS "Cancellation Rate"
FROM Trips
WHERE Request_at BETWEEN "2013-10-01" AND "2013-10-03"
AND Client_Id NOT IN (SELECT Users_Id FROM Users WHERE Banned = 'Yes')
AND Driver_Id NOT IN (SELECT Users_Id FROM Users WHERE Banned = 'Yes')
GROUP BY Request_at;
```

highest salary in the Sales department.

## Solution

SQL

```
SELECT Department.Name AS Department, Employee.Name AS Employee, Salary
FROM Employee
JOIN Department
ON Employee.DepartmentId = Department.Id
WHERE (DepartmentId, Salary) IN(
    SELECT DepartmentId, MAX(Salary) AS Salary
    FROM Employee
    GROUP BY DepartmentId
);
```

## 185. Department Top Th 6 | 0 | Hard | [LeetCode](#)

The Employee table holds all employees. Every employee has an Id, and there is also

The Employee table holds all employees. Every employee has an ID, and there is also a column for the department ID.

TEXT

Id   Name`   Salary   DepartmentId
1   Joe   85000   1
2   Henry   80000   2
3   Sam   60000   2
4   Max   90000   1
5   Janet   69000   1
6   Randy   85000   1
7   Will   70000   1

The Department table holds all departments of the company.

TEXT

Id	Name
1	IT
2	Sales

Write a SQL query to find employees who earn the top three salaries in each of the department. For the above tables, your SQL query should return the following rows (order of rows does not matter).

TEXT

Department	Employee	Salary
IT	Max	90000
IT	Randy	85000
IT	Joe	85000
IT	Will	70000
Sales	Henry	80000
Sales	Sam	60000

### ***Explanation:***

In IT department, Max earns the highest salary, both Randy and Joe earn the second highest salary, and Will earns the third highest salary. There are only two employees in the Sales department, Henry earns the highest salary while Sam earns the second highest salary.

 6 |  0

### **Solution**

SQL

```
WITH department_ranking AS (
SELECT Name AS Employee, Salary ,DepartmentId
,DENSE_RANK() OVER (PARTITION BY DepartmentId ORDER BY Salary DESC) AS rnk
FROM Employee
```



6 | 0

)  
SELECT d.Name AS Department, e.Employee, e.Salary  
FROM Employee e  
JOIN Department d ON e.DepartmentId = d.Id  
ORDER BY e.Salary DESC

```
SELECT d.Name AS Department, r.Employee, r.Salary
FROM department_ranking AS r
JOIN Department AS d
ON r.DepartmentId = d.Id
WHERE r.rnk <= 3
ORDER BY d.Name ASC, r.Salary DESC;
```

## 196. Delete Duplicate Emails | Easy | [LeetCode](#)

Write a SQL query to delete all duplicate email entries in a table named `Person`, keeping only unique emails based on its smallest Id.

TEXT

Id	Email
1	john@example.com
2	bob@example.com
3	john@example.com

`Id` is the primary key column for this table. For example, after running your query, the above `Person` table should have the following rows:

TEXT

Id	Email
1	john@example.com
2	bob@example.com

### Note:

Your output is the whole `Person` table after executing your sql. Use `delete`

statement.

## Solution

SQL

```
DELETE p2
FROM Person p1
JOIN Person p2
ON p1.Email = p2.Email
AND p1.id < p2.id
```



6 | 0

## Table: Weather

TEXT

Column Name	Type
id	int
recordDate	date
temperature	int

id is the primary key for this table.

This table contains information about the temperature in a certain day.

Write an SQL query to find all dates' id with higher temperature compared to its previous dates (yesterday).

Return the result table in any order.

The query result format is in the following example:

TEXT

Weather

id   recordDate   Temperature
1   2015-01-01   10
2   2015-01-02   25
3   2015-01-03   20
4   2015-01-04   30

Result table:

id
2
4

In 2015-01-02, temperature was higher than the previous day (10 -> 25).

In 2015-01-04, temperature was higher than the previous day (20 -> 30).

## Solution

SQL

6 | 0

#Solution- 1:

ANSWER



```
SELECT t.Id  
FROM Weather AS t, Weather AS y  
WHERE DATEDIFF(t.RecordDate, y.RecordDate) = 1  
AND t.Temperature > y.Temperature;
```

#Solution- 2:

```
SELECT t.Id  
FROM Weather t  
JOIN Weather y  
ON DATEDIFF(t.recordDate, y.recordDate) = 1 AND  
t.temperature > y.temperature;
```

## 262. Trips and Users | Hard | [LeetCode](#)

Table: Trips

TEXT

Column Name	Type
Id	int
Client_Id	int
Driver_Id	int
City_Id	int
Status	enum
Request_at	date

Id is the primary key for this table.

The table holds all taxi trips. Each trip has a unique Id, while Client\_Id and Driver\_Id are foreign keys to the Users table. Status is an ENUM type of ('completed', 'cancelled\_by\_driver', 'cancelled\_by\_client').

---

Table: Users

TEXT

Column Name	Type
Users_Id	int
Banned	enum
Role	enum

Users\_Id is the primary key for this table.

The table holds all users. Each user has a unique Users\_Id, and Role is an ENUM type of ('admin', 'user'). Banned is a boolean type.

---

Write a SQL query to find the cancellation rate of requests with unbanned users (both client and driver must not be banned) each day between "2013-10-01" and "2013-10-03".

The cancellation rate is computed by dividing the number of canceled (by client or driver) requests with unbanned users by the total number of requests with unbanned users on that day.

Return the result table in any order. Round cancellation Rate to two decimal points.

The query result format is in the following example:

TEXT

Trips table:

6 | 0

Id	Client_Id	Driver_Id	City_Id	Status	Request_at
1	1	10	1	completed	2013-10-01
2	2	11	1	cancelled_by_driver	2013-10-01
3	3	12	6	completed	2013-10-01
4	4	13	6	cancelled_by_client	2013-10-01
5	1	10	1	completed	2013-10-02
6	2	11	6	completed	2013-10-02
7	3	12	6	completed	2013-10-02
8	2	12	12	completed	2013-10-03
9	3	10	12	completed	2013-10-03
10	4	13	12	cancelled_by_driver	2013-10-03

Users table:

Users_Id	Banned	Role
1	No	client
2	Yes	client
3	No	client
4	No	client
10	No	driver
11	No	driver
12	No	driver
13	No	driver

Result table:

Day	Cancellation Rate
2013-10-01	0.33
2013-10-02	0.00
2013-10-03	0.50

On 2013-10-01:

- There were 4 requests in total, 2 of which were canceled.
- However, the request with Id=2 was made by a banned client (User\_Id=2), so it is ignored.
- Hence there are 3 unbanned requests in total, 1 of which was canceled.
- The Cancellation Rate is  $(1 / 3) = 0.33$

On 2013-10-02:

- There were 3 requests in total, 0 of which were canceled.
- The request with Id=6 was made by a banned client, so it is ignored.
- Hence there are 2 unbanned requests in total, 0 of which were canceled.
- The Cancellation Rate is  $(0 / 2) = 0.00$

On 2013-10-03:

- There were 3 requests in total, 1 of which was canceled.
- The request with Id=8 was made by a banned client, so it is ignored.
- Hence there are 2 unbanned requests in total, 1 of which were canceled.

- The Cancellation Rate is (1 / 2) = 0.50

---

## Solution

SQL

```
SELECT Request_at AS Day,  
ROUND(SUM(IF(Status<>"completed", 1, 0))/COUNT(Status),2) AS "Cancellation Rate"  
FROM Trips  
WHERE Request_at BETWEEN "2013-10-01" AND "2013-10-03"  
AND Client_Id NOT IN (SELECT Users_Id FROM Users WHERE Banned = 'Yes')  
AND Driver_Id NOT IN (SELECT Users_Id FROM Users WHERE Banned = 'Yes')  
GROUP BY Request_at;
```

---

## 511. Game Play Analysis I | Easy | 🔒 LeetCode

Table: Activity

TEXT

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player\_id, event\_date) is the primary key of this table.

This table shows the activity of players of some game.

Each row is a record of a player who logged in and played a number of games (poss

---

Write an SQL query that reports the **first login date** for each player.

The query result format is in the following example:

TEXT

Activity table:

player_id	device_id	event_date	games_played
1	2	2016-03-01	5
1	2	2016-05-02	6
2	3	2017-06-25	1
3	1	2016-03-02	0
3	4	2018-07-03	5

Like 6 | Dislike 0

Result table:

```
+-----+-----+
| player_id | first_login |
+-----+-----+
| 1          | 2016-03-01   |
```

```
+-----+-----+
| 2          | 2017-06-25   |
| 3          | 2016-03-02   |
+-----+-----+
```

赞同 6 | 反对 0

## Solution

SQL

```
SELECT player_id, MIN(event_date) as first_login
FROM Activity
GROUP BY player_id
```



## 512. Game Play Analysis II | Easy | [LeetCode](#)

Table: Activity

TEXT

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player\_id, event\_date) is the primary key of this table.

This table shows the activity of players of some game.

Each row is a record of a player who logged in and played a number of games (poss



Write a SQL query that reports the device that is first logged in for each player.

The query result format is in the following example:

TEXT

Activity table:

player_id	device_id	event_date	games_played
1	2	2016-03-01	5
1	2	2016-05-02	6
2	3	2017-06-25	1
3	1	2016-03-02	0
3	4	2018-07-03	5

Result table:

player_id	device_id	upvotes	downvotes
1	2	6	0
2	3		

3	1	
+-----+-----+		

## Solution

SQL

```
#Solution- 1:  
SELECT DISTINCT player_id, device_id  
FROM Activity  
WHERE (player_id, event_date) in (  
    SELECT player_id, min(event_date)  
    FROM Activity  
    GROUP BY player_id)
```

```
#Solution- 2:
```

```
SELECT a.player_id, b.device_id  
FROM  
(SELECT player_id, MIN(event_date) AS event_date FROM Activity  
GROUP BY player_id) a  
JOIN Activity b
```

```
ON a.player_id = b.player_id AND a.event_date = b.event_date;
```

#Solution- 3:

```
SELECT player_id, device_id
FROM
(SELECT player_id, device_id, event_date,
ROW_NUMBER() OVER (PARTITION BY player_id ORDER BY event_date) AS r
FROM Activity) lookup
WHERE r = 1;
```

## 534. Game Play Analysis III | Medium | [LeetCode](#)

Table: Activity

TEXT

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player\_id, event\_date) is the primary key of this table.

This table shows the activity of players of some game.

Each row is a record of a player who logged in and played a number of games (poss

 6 |  0

TEXT

id	name	referee_id
1	Will	NULL
2	Jane	NULL
3	Alex	2
4	Bill	NULL
5	Zack	1
6	Mark	2

Write a query to return the list of customers NOT referred by the person with id '2'.

For the sample data above, the result is:

TEXT

name
Will
Jane
Bill
Zack

## Solution

SQL

```
SELECT name
FROM customer
WHERE referee_id != '2' OR referee_id IS NULL;
```

**585. Investments in 2016 | Medium | 🔒 LeetCode**

Write a query to print the sum of all total investment values in 2016 (TIV\_2016), to a scale of 2 decimal places, for all policy holders who meet the following criteria:

1. Have the same TIV\_2015 value as one or more other policyholders.
2. Are not located in the same city as any other policyholder (i.e.: the (latitude, longitude) attribute pairs must be unique). Input Format: The insurance table is described as follows:

TEXT

Column Name	Type
PID	INTEGER(11)
TIV_2015	NUMERIC(15,2)
TIV_2016	NUMERIC(15,2)
LAT	NUMERIC(5,2)
LON	NUMERIC(5,2)

where PID is the policyholder's policy ID, TIV2015 is the total investment value in 2015, TIV2016 is the total investment value in 2016, LAT is the latitude of the policy holder's city, and LON is the longitude of the policy holder's city.

#### Sample Input

TEXT

PID	TIV_2015	TIV_2016	LAT	LON
1	10	5	10	10
2	20	20	20	20
3	10	30	20	20
4	10	40	40	40

#### Sample Output

TEXT

TIV_2016
45.00

## Explanation

TEXT

The first record in the table, like the last record, meets both of the two criteria. The TIV\_2015 value '10' is as the same as the third and forth record, and its location is the same.

The second record does not meet any of the two criteria. Its TIV\_2015 is not like 10, and its location is not the same with the first record.

And its location is the same with the third record, which makes the third record meet both criteria.

So, the result is the sum of TIV\_2016 of the first and last record, which is 45.

---

## Solution

SQL

```
SELECT SUM(TIV_2016) AS TIV_2016
FROM insurance
WHERE CONCAT(LAT, ',', LON)
    IN (SELECT CONCAT(LAT, ',', LON)
        FROM insurance
        GROUP BY LAT, LON
        HAVING COUNT(1) = 1)
AND TIV_2015 in
    (SELECT TIV_2015
    FROM insurance
    GROUP BY TIV_2015
    HAVING COUNT(1)>1)
```



## 586. Customer Placing the Largest Number of Orders | Easy | LeetCode

Query the customer\_number from the orders table for the customer who has placed the largest number of orders.

6 | 0

It is guaranteed that exactly one customer will have placed more orders than any other customer.

The orders table is defined as follows:

TEXT

Column	Type
order_number (PK)	int
customer_number	int
order_date	date
required_date	date
shipped_date	date
status	char(15)
comment	char(200)

### Sample Input

TEXT

order_number	customer_number	order_date	required_date	shipped_date	st
1	1	2017-04-09	2017-04-13	2017-04-12	C1
2	2	2017-04-15	2017-04-20	2017-04-18	C1
3	3	2017-04-16	2017-04-25	2017-04-20	C1
4	3	2017-04-18	2017-04-28	2017-04-25	C1

=====  
=====

### Sample Output

TEXT

customer_number
3

### Explanation

6 | 0

TEXT

The customer with number '3' has two orders, which is greater than either customer with number '1' or '2'. So the result is customer\_number '3'.

---

## Solution

SQL

```
# assume: only one match
SELECT customer_number FROM orders
GROUP BY customer_number
ORDER BY COUNT(1) DESC
LIMIT 1

## assume: multiple matches
## 1 1
## 2 1
## 3 1
##
## 1 1 1 1
## 1 1 2 1
## 1 1 3 1
##
## SELECT t1.customer_number
##   FROM (SELECT customer_number, COUNT(1) AS count
##          FROM orders GROUP BY customer_number) AS t1,
##       (SELECT customer_number, COUNT(1) AS count
##          FROM orders GROUP BY customer_number) AS t2
##   GROUP BY t1.customer_number
##   HAVING max(t1.count) = max(t2.count)
```



## 595. Big Countries | Easy | [LeetCode](#)

There is a table World

TEXT

name	continent	area	population	gdp
Afghanistan	Asia	652230	25500100	20343000
Albania	Europe	28748	2831741	12960000
Algeria	Africa	2381741	37100000	188681000
Andorra	Europe	468	78115	3712000
Angola	Africa	1246700	20609294	100990000

A country is big if it has an area of bigger than 3 million square km or a population of more than 25 million.

Write a SQL solution to output big countries' name, population and area.

For example, according to the above table, we should output:

TEXT

name	population	area
Afghanistan	25500100	652230
Algeria	37100000	2381741

## Solution

SQL

```
SELECT name, population, area
FROM World
WHERE area >= 3000000 OR population > 25000000;
```



## 596. Classes More Than 5 Students | Easy | [LeetCode](#)

6 | 0

There is a table `courses` with columns: **student** and **class**

Please list out all classes which have more than or equal to 5 students.

For example, the table:

TEXT	
+-----+-----+	
student	class
+-----+-----+	
A	Math
B	English
C	Math
D	Biology
E	Math
F	Computer
G	Math
H	Math
I	Math
+-----+-----+	

Should output:

TEXT	
+-----+	
class	
+-----+	
Math	
+-----+	

## Solution

SQL

```
SELECT class
FROM courses
GROUP BY class
HAVING count(DISTINCT Student)>=5;
```



## 597. Friend Requests I: Overall Acceptance Rate | Easy |

[LeetCode](#)

In social network like Facebook or Twitter, people send friend requests and accept others' requests as well. Now given two tables as below: Table: friend\_request

TEXT

sender_id	send_to_id	request_date
1	2	2016_06-01
1	3	2016_06-01
1	4	2016_06-01
2	3	2016_06-02
3	4	2016-06-09

Table: request\_accepted

TEXT

requester_id	accepter_id	accept_date
1	2	2016_06-03
1	3	2016-06-08
2	3	2016-06-08
3	4	2016-06-09
3	4	2016-06-10

Write a query to find the overall acceptance rate of requests rounded to 2 decimals, which is the number of acceptance divide the number of requests. For the sample data above, your query should return the following result.

TEXT

accept_rate
0.80

Note:

The accepted requests are not necessarily from the table friendrequest. In this case, you just need to simply count the total accepted requests (no matter whether they are in the original requests), and divide it by the number of requests to get the acceptance rate. It is possible that a sender sends multiple requests to the same receiver, and a request could be accepted more than once. In this case, the 'duplicated' requests or acceptances are only counted once. If there is no requests at all, you should return 0.00 as the acceptrate. Explanation: There are 4 unique accepted requests, and there are 5 requests in total. So the rate is 0.80.

Follow-up:

Can you write a query to return the accept rate but for every month? How about the cumulative accept rate for every day?

## Solution

SQL

```
SELECT IFNULL((round(accepts/requests, 2)), 0.0) AS accept_rate
FROM
    (SELECT count(DISTINCT sender_id, send_to_id) AS requests FROM friend_request
     (SELECT count(DISTINCT requester_id, accepter_id) AS accepts FROM request_acc
```

## 601. Human Traffic of Stadium | Hard | [LeetCode](#)

Table: stadium

TEXT

Column Name	Type
id	int
visit_date	date
people	int

*visitdate is the primary key for this table. Each row of this table contains the visit date and visit id to the stadium with the number of people during the visit. No two rows will have the same visitdate, and as the id increases, the dates increase as well.*

Write an SQL query to display the records with three or more rows with **consecutive** id's, and the number of people is greater than or equal to 100 for each.

Return the result table ordered by `visit_date` in **ascending order**.

The query result format is in the following example.

TEXT

Stadium table:

id	visit_date	people
1	2017-01-01	10
2	2017-01-02	109
3	2017-01-03	150
4	2017-01-04	99
5	2017-01-05	145
6	2017-01-06	1455
7	2017-01-07	199
8	2017-01-09	188

Result table:

id	visit_date	people
5	2017-01-05	145
6	2017-01-06	1455
7	2017-01-07	199
8	2017-01-09	188

The four rows with ids 5, 6, 7, and 8 have consecutive ids and each of them has > The rows with ids 2 and 3 are not included because we need at least three consecu

## Solution

```
SQL
SELECT DISTINCT s1.*
FROM Stadium s1 JOIN Stadium s2 JOIN Stadium s3
ON (s1.id = s2.id-1 AND s1.id = s3.id-2) OR
(s1.id = s2.id+1 AND s1.id = s3.id-1) OR
(s1.id = s2.id+1 AND s1.id = s3.id+2)
WHERE s1.people >= 100 AND s2.people >= 100 AND s3.people>=100
ORDER BY visit_date
```



## 602. Friend Requests II: Who Has the Most Friends | Medium |



[LeetCode](#)

In social network like Facebook or Twitter, people send friend requests and accept others' requests as well. Table `request_accepted` holds the data of friend acceptance, while `requesterid` and `accepterid` both are the id of a person.

TEXT

requester_id	accepter_id	accept_date
1	2	2016_06-03
1	3	2016-06-08
2	3	2016-06-08
3	4	2016-06-09

Write a query to find the the people who has most friends and the most friends number. For the sample data above, the result is:

TEXT

id	num
3	3

6 | 0

Note:

It is guaranteed there is only 1 people having the most friends. The friend request could only been accepted once, which mean there is no multiple records with the same requesterid and accepterid value. Explanation: The person with id '3' is a friend of people '1', '2' and '4', so he has 3 friends in total, which is the most number than any others.

Follow-up: In the real world, multiple people could have the same most number of friends, can you find all these people in this case?

SQL

```
SELECT t.id, sum(t.num) AS num
FROM (
    (SELECT requester_id AS id, COUNT(1) AS num
     FROM request_accepted
     GROUP BY requester_id)
    union all
    (SELECT accepter_id AS id, COUNT(1) AS num
     FROM request_accepted
     GROUP BY accepter_id)) AS t
GROUP BY t.id
ORDER BY num DESC
LIMIT 1;
```



## 603. Consecutive Available Seats | Easy | 🔒 LeetCode

Several friends at a cinema ticket office would like to reserve consecutive available seats. Can you help to query all the consecutive available seats order by the seat\_id using the following cinema table?

TEXT

seat_id	free
1	1
2	0
3	1

4	1	
5	1	

Your query should return the following result for the sample case above.

TEXT
seat_id
-----
3
4
5

Note:

The seat\_id is an auto increment int, and free is bool ('1' means free, and '0' means occupied.). Consecutive available seats are more than 2(inclusive) seats consecutively available.

## Solution

SQL

```
SELECT DISTINCT t1.seat_id
FROM cinema AS t1 JOIN cinema AS t2
ON abs(t1.seat_id-t2.seat_id)=1
WHERE t1.free='1' AND t2.free='1'
ORDER BY t1.seat_id
```



## 607.Sales Person | Easy | 🔒 [LeetCode](#)

### Description

Given three tables: `salesperson`, `company`, `orders`. Output all the names in the table `salesperson`, who didn't have sales to company 'RED'.

### Example Input

6 | 0

Table: salesperson

TEXT

sales_id	name	salary	commission_rate	hire_date
1	John	100000	6	4/1/2006
2	Amy	120000	5	5/1/2010
3	Mark	65000	12	12/25/2008
4	Pam	25000	25	1/1/2005
5	Alex	50000	10	2/3/2007

The table salesperson holds the salesperson information. Every salesperson has a sales\_id and a name. Table: company

TEXT

com_id	name	city
1	RED	Boston
2	ORANGE	New York
3	YELLOW	Boston
4	GREEN	Austin

The table company holds the company information. Every company has a com\_id and a name. Table: orders

TEXT

order_id	date	com_id	sales_id	amount
1	1/1/2014	3	4	100000
2	2/1/2014	4	5	5000
3	3/1/2014	1	1	50000
4	4/1/2014	1	4	25000

The table orders holds the sales record information, salesperson and customer company are represented by salesid and comid. output

```
TEXT
+-----+
| name |
+-----+
| Amy   |
| Mark  |
| Alex  |
+-----+
```

## Explanation

According to order '3' and '4' in table orders, it is easy to tell only salesperson 'John' and 'Alex' have sales to company 'RED', so we need to output all the other names in table salesperson.

## Solution

```
SQL
SELECT name
FROM salesperson
WHERE name NOT IN
    (SELECT DISTINCT salesperson.name
     FROM salesperson, orders, company
     WHERE company.name = 'RED'
     AND salesperson.sales_id = orders.sales_id
     AND orders.com_id = company.com_id)
```

## 608. Tree Node | Medium | [LeetCode](#)

Given a table tree, id is identifier of the tree node and p\_id is its parent node's id.

```
TEXT
+-----+
| id | p_id |
+-----+
| 1  | null  |
| 2  | 1      |
| 3  | 1      |
| 4  | 2      |
| 5  | 2      |
+-----+
```

Each node in the tree can be one of three types:

Leaf: if the node is a leaf node. Root: if the node is the root of the tree. Inner: If the node is neither a leaf node nor a root node. Write a query to print the node id and the type of the node. Sort your output by the node id. The result for the above sample is:

```
TEXT
+-----+
| id | Type |
+-----+
| 1  | Root  |
| 2  | Inner |
| 3  | Leaf   |
| 4  | Leaf   |
| 5  | Leaf   |
+-----+
```

### Explanation

Node '1' is root node, because its parent node is NULL and it has child node '2' and '3'. Node '2' is inner node, because it has parent node '1' and child node '4' and '5'. Node '3', '4' and '5' is Leaf node, because they have parent node and they don't have child node. And here is the image of the sample tree as below:

TEXT

```
1
 /   \
2     3
 /   \
4     5
```

## Note

If there is only one node on the tree, you only need to output its root attributes.

## Solution

SQL

```
## Basic Idea: LEFT JOIN
# In tree, each node can only one parent or no parent
## | id | p_id | id (child) |
## |-----+-----|
## | 1 | null |          1 |
## | 1 | null |          2 |
## | 2 | ... 1 |          4 |
## | 2 | ... 1 |          5 |
## | 3 | ... 1 |          null |
## | 4 | ... 2 |          null |
## | 5 | ... 2 |          null |

SELECT t1.id,
CASE
    WHEN ISNULL(t1.p_id) THEN 'Root'
    WHEN ISNULL(MAX(t2.id)) THEN 'Leaf'
    ELSE 'Inner'
END AS Type
FROM tree AS t1 LEFT JOIN tree AS t2
ON t1.id = t2.p_id
GROUP BY t1.id, t1.p_id
```

## 610. Triangle Judgement | Easy | 🔒 LeetCode

A pupil Tim gets homework to identify whether three line segments could possibly form a triangle. However, this assignment is very heavy because there are hundreds of records to calculate. Could you help Tim by writing a query to judge whether these three sides can form a triangle, assuming table triangle holds the length of the three sides x, y and z.

TEXT

x	y	z
13	15	30
10	20	15

For the sample data above, your query should return the follow result:

TEXT

x	y	z	triangle
13	15	30	No
10	20	15	Yes

## Solution

SQL

```
SELECT x, y, z,
CASE
    WHEN x+y>z AND y+z>x AND x+z>y THEN 'Yes'
    ELSE 'No'
END AS triangle
FROM triangle
```



## 612. Shortest Distance in a Plane | Medium | 🔒 LeetCode

Table point\_2d holds the coordinates (x,y) of some unique points (more than two) in a plane. Write a query to find the shortest distance between these points rounded to 2 decimals.

TEXT

x	y
-1	-1
0	0
-1	-2

The shortest distance is 1.00 from point (-1,-1) to (-1,2). So the output should be:

TEXT

shortest
-----
1.00

Note: The longest distance among all the points are less than 10000.

## Solution

SQL

```
SELECT ROUND(MIN(SQRT((t1.x-t2.x)*(t1.x-t2.x) + (t1.y-t2.y)*(t1.y-t2.y))), 2) AS shortest
FROM point_2d AS t1, point_2d AS t2
WHERE t1.x!=t2.x OR t1.y!=t2.y

# SELECT ROUND(SQRT((t1.x-t2.x)*(t1.x-t2.x) + (t1.y-t2.y)*(t1.y-t2.y)), 2) AS shortest
# FROM point_2d AS t1, point_2d AS t2
# WHERE t1.x!=t2.x OR t1.y!=t2.y
# ORDER BY shortest ASC
# LIMIT 1
```

**613. Shortest Distance in a Line | Easy |  [LeetCode](#)**

Table point holds the x coordinate of some points on x-axis in a plane, which are all integers. Write a query to find the shortest distance between two points in these points.

TEXT
x
-----
-1
0
2

The shortest distance is '1' obviously, which is from point '-1' to '0'. So the output is as below:

TEXT
shortest
-----
1

Note: Every point is unique, which means there is no duplicates in table point.

Follow-up: What if all these points have an id and are arranged from the left most to the right most of x axis?

## Solution

```
SQL
SELECT t1.x-t2.x AS shortest
FROM point AS t1 JOIN point AS t2
WHERE t1.x>t2.x
ORDER BY (t1.x-t2.x) ASC
LIMIT 1
```



**614. Second Degree Follower | Medium | 🔒 LeetCode**

In facebook, there is a follow table with two columns: followee, follower.

Please write a sql query to get the amount of each follower's follower if he/she has one.

For example:

TEXT		
followee	follower	
A	B	
B	C	
B	D	
D	E	

should output:

TEXT		
follower	num	
B	2	
D	1	

Explanation: Both B and D exist in the follower list, when as a followee, B's follower is C and D, and D's follower is E. A does not exist in follower list.

Note: Followee would not follow himself/herself in all cases. Please display the result in follower's alphabet order.

## Solution



```

SQL
## Explain the business logic
## . A follows B. Then A is follower, B is followee
## What are second degree followers?
## . A follows B, and B follows C.
## . Then A is the second degree followers of C

SELECT f1.follower, COUNT(DISTINCT f2.follower) AS num
FROM follow AS f1 JOIN follow AS f2
ON f1.follower = f2.followee
GROUP BY f1.follower;

```

## 615. Average Salary: Departments VS Company | Hard |

[LeetCode](#)

Given two tables as below, write a query to display the comparison result (higher/lower/same) of the average salary of employees in a department to the company's average salary. Table: salary

TEXT

id	employee_id	amount	pay_date
1	1	9000	2017-03-31
2	2	6000	2017-03-31
3	3	10000	2017-03-31
4	1	7000	2017-02-28
5	2	6000	2017-02-28
6	3	8000	2017-02-28

The *employeeid* column refers to the *employeeid* in the following table employee.

TEXT

employee_id	department_id
1	1
2	2
3	2

So for the sample data above, the result is:

pay_month	department_id	comparison
2017-03	1	higher
2017-03	2	lower
2017-02	1	same
2017-02	2	same

Explanation In March, the company's average salary is  $(9000+6000+10000)/3 = 8333.33$ ... The average salary for department '1' is 9000, which is the salary of employeeid '1' since there is only one employee in this department. So the comparison result is 'higher' since  $9000 > 8333.33$  obviously. The average salary of department '2' is  $(6000 + 10000)/2 = 8000$ , which is the average of employeeid '2' and '3'. So the comparison result is 'lower' since  $8000 < 8333.33$ . With the same formula for the average salary comparison in February, the result is 'same' since both the department '1' and '2' have the same average salary with the company, which is 7000.

## Solution

```
SQL   
SELECT t1.pay_month, t1.department_id,  
       (CASE WHEN t1.amount = t2.amount THEN 'same'  
             WHEN t1.amount > t2.amount THEN 'higher'  
             WHEN t1.amount < t2.amount THEN 'lower' END) AS comparison  
FROM  
       (SELECT left(pay_date, 7) AS pay_month, department_id, avg(amount) AS amount  
        FROM salary JOIN employee  
        ON salary.employee_id = employee.employee_id  
        GROUP BY pay_month, department_id  
        ORDER BY pay_month DESC, department_id) AS t1  
       JOIN  
       (SELECT left(pay_date, 7) AS pay_month, avg(amount) AS amount  
        FROM salary JOIN employee
```

```
ON salary.employee_id = employee.employee_id  
GROUP BY pay_month) AS t2  
ON t1.pay_month = t2.pay_month
```

□

## 618. Students Report By Geography | Hard | 🔒 LeetCode

A U.S graduate school has students from Asia, Europe and America. The students' location information are stored in table student as below.

TEXT

name	continent
Jack	America
Pascal	Europe
Xi	Asia
Jane	America

Pivot the continent column in this table so that each name is sorted alphabetically and displayed underneath its corresponding continent. The output headers should be America, Asia and Europe respectively. It is guaranteed that the student number from America is no less than either Asia or Europe. For the sample input, the output is:

TEXT

America	Asia	Europe
Jack	Xi	Pascal
Jane		

Follow-up: If it is unknown which continent has the most students, can you write a query to generate the student report?

### Solution

ANS

1 6 | 0

SQL

```
SELECT t1.name AS America, t2.name AS Asia, t3.name AS Europe
FROM
  (SELECT (@cnt1 := @cnt1 + 1) AS id, name
   FROM student
   CROSS JOIN (SELECT @cnt1 := 0) AS dummy
   WHERE continent='America'
   ORDER BY name) AS t1
  LEFT JOIN
  (SELECT (@cnt2 := @cnt2 + 1) AS id, name
   FROM student
   CROSS JOIN (SELECT @cnt2 := 0) AS dummy
   WHERE continent='Asia'
   ORDER BY name) AS t2
  ON t1.id = t2.id
  LEFT JOIN
  (SELECT (@cnt3 := @cnt3 + 1) AS id, name
   FROM student
   CROSS JOIN (SELECT @cnt3 := 0) AS dummy
   WHERE continent='Europe'
   ORDER BY name) AS t3
  ON t1.id = t3.id
```

## 619. Biggest Single Number | Easy | 🔒 [LeetCode](#)

Table number contains many numbers in column num including duplicated ones. Can you write a SQL query to find the biggest number, which only appears once.

TEXT

num
8
8
3
3
1
4

```
| 5 |
| 6 |
```

For the sample data above, your query should return the following result:

```
TEXT
+---+
| num |
+---+
| 6 |
```

Note: If there is no such number, just output null.

## Solution

```
SQL
SELECT IFNULL(
    SELECT num
    FROM number
    GROUP BY num
    HAVING count(1) = 1
    ORDER BY num DESC
    LIMIT 0, 1), NULL) AS num
```

## 620. Not Boring Movies | Easy | [LeetCode](#)

X city opened a new cinema, many people would like to go to this cinema. The cinema also gives out a poster indicating the movies' ratings and descriptions. Please write a SQL query to output movies with an odd numbered ID and a description that is not 'boring'. Order the result by rating.

For example, table `cinema`:

TEXT

id	movie	description	rating
1	War	great 3D	8.9
2	Science	fiction	8.5
3	irish	boring	6.2
4	Ice song	Fantacy	8.6
5	House card	Interesting	9.1

For the example above, the output should be:

TEXT

id	movie	description	rating
5	House card	Interesting	9.1
1	War	great 3D	8.9

## Solution

SQL

```
SELECT *
FROM Cinema
WHERE description <> 'boring' AND ID % 2 = 1
ORDER BY rating DESC;
```



## 626. Exchange Seats | Medium | [LeetCode](#)

Mary is a teacher in a middle school and she has a table `seat` storing students' names and their corresponding seat ids.

The column `id` is continuous increment.

Mary wants to change seats for the adjacent students.

Can you write a SQL query to output the result for Mary?

TEXT

id	student
1	Abbot
2	Doris
3	Emerson
4	Green
5	Jeames

For the sample input, the output is:

TEXT

id	student
1	Doris
2	Abbot
3	Green
4	Emerson
5	Jeames

### Note:

If the number of students is odd, there is no need to change the last one's seat.

## Solution

SQL

```
SELECT  
IF(id<(SELECT MAX(id) FROM seat),IF(id%2=0,id-1, id+1),IF(id%2=0, id-1, id)) AS i
```

```
FROM seat  
ORDER BY id;
```

=====

## 627. Swap Salary | [LeetCode](#)

Table: salary

TEXT

Column Name	Type
id	int
name	varchar
sex	ENUM
salary	int

id is the primary key for this table.

The sex column is ENUM value of type ('m', 'f').

The table contains information about an employee.

Write an SQL query to swap all 'f' and 'm' values (i.e., change all 'f' values to 'm' and vice versa) with a single update statement and no intermediate temp table(s).

Note that you must write a single update statement, DO NOT write any select statement for this problem.

The query result format is in the following example:

TEXT

Salary table:

id	name	sex	salary
1	A	m	2500

```

| 2 | B   | f   | 1500 |
| 3 | C   | m   | 5500 |
| 4 | D   | f   | 500  |
+-----+-----+

```

Result table:

```

+-----+-----+-----+
| id | name | sex | salary |
+-----+-----+-----+
| 1  | A    | f   | 2500 |
| 2  | B    | m   | 1500 |
| 3  | C    | f   | 5500 |
| 4  | D    | m   | 500  |
+-----+-----+-----+

```

(1, A) and (2, C) were changed from 'm' to 'f'.

(2, B) and (4, D) were changed from 'f' to 'm'.

## Solution

SQL



# With IF

```
UPDATE Salary SET sex = IF(sex='m', 'f', 'm')
```

# With CASE

```
UPDATE Salary SET sex = CASE WHEN sex='m' THEN 'f' ELSE 'm' END
```

## 1045. Customers Who Bought All Products | Medium | 🔒 LeetCode

Table: Customer

TEXT

```

+-----+-----+
| Column Name | Type   |
+-----+-----+
| customer_id | int   |

```

```
| product_key | int      |  
+-----+-----+
```

product\_key is a foreign key to Product table. Table: Product

TEXT

```
+-----+-----+  
| Column Name | Type      |  
+-----+-----+  
| product_key | int      |  
+-----+-----+  
product_key is the primary key column for this table.
```

Write an SQL query for a report that provides the customer ids from the Customer table that bought all the products in the Product table.

For example:

TEXT

Customer table:

```
+-----+-----+  
| customer_id | product_key |  
+-----+-----+  
| 1           | 5          |  
| 2           | 6          |  
| 3           | 5          |  
| 3           | 6          |  
| 1           | 6          |  
+-----+-----+
```

Product table:

```
+-----+  
| product_key |  
+-----+  
| 5           |  
| 6           |  
+-----+
```

Result table:

customer_id
1
3

The customers who bought all the products (5 and 6) are customers with id 1 and 3



## Solution

SQL



```
SELECT customer_id
FROM Customer
GROUP BY customer_id
HAVING count(DISTINCT product_key) = (
    SELECT count(1)
    FROM Product)
```

## 1050. Actors and Directors Who Cooperated At Least Three Times | Easy | [LeetCode](#)

Table: ActorDirector

TEXT

Column Name	Type
actor_id	int
director_id	int
timestamp	int

timestamp is the primary key column for this table.

Write a SQL query for a report that provides the pairs (*actorid*, *directorid*) where the actor have cooperated with the director at least 3 times.

Example:

TEXT

ActorDirector table:

actor_id	director_id	timestamp
1	1	0
1	1	1
1	1	2
1	2	3
1	2	4
2	1	5
2	1	6

Result table:

actor_id	director_id
1	1

The only pair is (1, 1) where they cooperated exactly 3 times.

## Solution

SQL

```
SELECT actor_id, director_id
FROM ActorDirector
GROUP BY actor_id, director_id
HAVING COUNT(1)>=3
```



## 1068. Product Sales Analysis I | Easy | 🔒 LeetCode

1 6 | 0

### Table: Sales

TEXT

Column Name	Type
sale_id	int
product_id	int
year	int
quantity	int
price	int

(sale\_id, year) is the primary key of this table.  
product\_id is a foreign key to Product table.  
Note that the price is per unit.

### Table: Product

TEXT

Column Name	Type
product_id	int
product_name	varchar

product\_id is the primary key of this table.

Write an SQL query that reports all product names of the products in the Sales table along with their selling year and price.

For example:

TEXT

Sales table:

sale_id	product_id	year	quantity	price
1	100	2008	10	5000

2	100	2009	12	5000
7	200	2011	15	9000

Product table:

product_id	product_name
100	Nokia
200	Apple
300	Samsung

Result table:

product_name	year	price
Nokia	2008	5000
Nokia	2009	5000
Apple	2011	9000

## Solution

SQL

```
SELECT product_name, year, price
FROM Sales JOIN Product
ON Product.product_id = Sales.product_id
```



## 1069. Product Sales Analysis II | Easy | 🔒 LeetCode

Table: Sales

TEXT

Column Name	Type
-------------	------

```

+-----+-----+
| sale_id | int   |
| product_id | int   |
| year     | int   |
| quantity | int   |
| price    | int   |
+-----+-----+
sale_id is the primary key of this table.
product_id is a foreign key to Product table.
Note that the price is per unit.

```

### Table: Product

```

TEXT
+-----+-----+
| Column Name | Type   |
+-----+-----+
| product_id  | int    |
| product_name | varchar |
+-----+-----+
product_id is the primary key of this table.

```

Write an SQL query that reports the total quantity sold for every product id.

The query result format is in the following example:

```

TEXT
Sales table:
+-----+-----+-----+-----+
| sale_id | product_id | year | quantity | price |
+-----+-----+-----+-----+
| 1       | 100        | 2008 | 10      | 5000  |
| 2       | 100        | 2009 | 12      | 5000  |
| 7       | 200        | 2011 | 15      | 9000  |
+-----+-----+-----+-----+
Product table:
+-----+-----+

```

```

| product_id | product_name |
+-----+-----+
| 100      | Nokia      |
| 200      | Apple       |
| 300      | Samsung    |
+-----+-----+

```

Result table:

```

+-----+-----+
| product_id | total_quantity |
+-----+-----+
| 100      | 22          |
| 200      | 15          |
+-----+-----+

```

## Solution

SQL

```

SELECT product_id, sum(quantity) AS total_quantity
FROM Sales
GROUP BY product_id;

```



## 1070. Product Sales Analysis III | Medium | 🔒 LeetCode

Table: sales

TEXT

```

+-----+-----+
| Column Name | Type   |
+-----+-----+
| sale_id     | int    |
| product_id  | int    |
| year        | int    |
| quantity    | int    |
| price        | int    |
+-----+-----+

```

sale\_id is the primary key of this table.  
product\_id is a foreign key to Product table.  
Note that the price is per unit.

#### Table: Product

TEXT

Column Name	Type
product_id	int
product_name	varchar

product\_id is the primary key of this table.

Write an SQL query that selects the product id, year, quantity, and price for the first year of every product sold.

The query result format is in the following example:

TEXT

Sales table:

sale_id	product_id	year	quantity	price
1	100	2008	10	5000
2	100	2009	12	5000
7	200	2011	15	9000

Product table:

product_id	product_name
100	Nokia
200	Apple
300	Samsung

Result table:

product_id	first_year	quantity	price
100	2008	10	5000
200	2011	15	9000

## Solution

SQL



```
SELECT
    product_id,
    year first_year,
    quantity,
    price
FROM Sales
WHERE (product_id, year) IN (SELECT product_id, MIN(year)
                               FROM Sales
                               GROUP BY product_id)
```

## 1075. Project Employees I | Easy | [LeetCode](#)

Table: Project

TEXT

Column Name	Type
project_id	int
employee_id	int

(project\_id, employee\_id) is the primary key of this table.  
employee\_id is a foreign key to Employee table.

### Table: Employee

```
TEXT
+-----+-----+
| Column Name      | Type     |
+-----+-----+
| employee_id      | int      |
| name              | varchar  |
| experience_years | int      |
+-----+-----+
employee_id is the primary key of this table.
```

Write an SQL query that reports the average experience years of all the employees for each project, rounded to 2 digits.

The query result format is in the following example:

```
TEXT
Project table:
+-----+-----+
| project_id | employee_id |
+-----+-----+
| 1          | 1           |
| 1          | 2           |
| 1          | 3           |
| 2          | 1           |
| 2          | 4           |
+-----+-----+

Employee table:
+-----+-----+-----+
| employee_id | name    | experience_years |
+-----+-----+-----+
| 1           | Khaled  | 3             |
| 2           | Ali     | 2             |
| 3           | John    | 1             |
| 4           | Doe    | 2             |
+-----+-----+
```

## 597. Friend Requests I: Overall Acceptance Rate | Easy |

### LeetCode

In social network like Facebook or Twitter, people send friend requests and accept others' requests as well. Now given two tables as below: Table: friend\_request

TEXT

sender_id	send_to_id	request_date
1	2	2016_06-01
1	3	2016_06-01
1	4	2016_06-01
2	3	2016_06-02
3	4	2016-06-09

Table: request\_accepted

TEXT

requester_id	accepter_id	accept_date
1	2	2016_06-03
1	3	2016-06-08
2	3	2016-06-08
3	4	2016-06-09
3	4	2016-06-10

 6 |  0

Write a query to find the overall acceptance rate of requests rounded to 2 decimals,

which is the number of acceptance divide the number of requests. For the sample data above, your query should return the following result.

TEXT

accept_rate
-----
0.80

6 | 0

Note:

The accepted requests are not necessarily from the table friendrequest. In this case, you just need to simply count the total accepted requests (no matter whether they are in the original requests), and divide it by the number of requests to get the acceptance rate. It is possible that a sender sends multiple requests to the same receiver, and a request could be accepted more than once. In this case, the 'duplicated' requests or acceptances are only counted once. If there is no requests at all, you should return 0.00 as the acceptrate. Explanation: There are 4 unique accepted requests, and there are 5 requests in total. So the rate is 0.80.

Follow-up:

Can you write a query to return the accept rate but for every month? How about the cumulative accept rate for every day?

## Solution

```
SQL
SELECT IFNULL((round(accepts/requests, 2)), 0.0) AS accept_rate
FROM
  (SELECT count(DISTINCT sender_id, send_to_id) AS requests FROM friend_request
  (SELECT count(DISTINCT requester_id, accepter_id) AS accepts FROM request_acc
```

## 601. Human Traffic of Stadium | Hard | [LeetCode](#)

Table: stadium

```
TEXT
+-----+-----+
| Column Name | Type   |
+-----+-----+
| id          | int    |
| visit_date  | date   |
| people      | int    |
+-----+-----+
```

*visitdate is the primary key for this table. Each row of this table contains the visit date and visit id to the stadium with the number of people during the visit. No two rows will have the same visitdate, and as the id increases, the dates increase as well.*

Write an SQL query to display the records with three or more rows with **consecutive** id's, and the number of people is greater than or equal to 100 for each.

Return the result table ordered by `visit_date` in **ascending order**.

The query result format is in the **following example**.

 6 |  0

Stadium table:

id	visit_date	people	
1	2017-01-01	10	
2	2017-01-02	109	
3	2017-01-03	150	
4	2017-01-04	99	
5	2017-01-05	145	
6	2017-01-06	1455	
7	2017-01-07	199	
8	2017-01-09	188	

Result table:

id	visit_date	people	
5	2017-01-05	145	
6	2017-01-06	1455	
7	2017-01-07	199	
8	2017-01-09	188	

The four rows with ids 5, 6, 7, and 8 have consecutive ids and each of them has > 100 people. The rows with ids 2 and 3 are not included because we need at least three consecutive rows.

---

## Solution

SQL

```
SELECT DISTINCT s1.*  
FROM Stadium s1 JOIN Stadium s2 JOIN Stadium s3  
ON (s1.id = s2.id-1 AND s1.id = s3.id-2) OR  
(s1.id = s2.id+1 AND s1.id = s3.id-1) OR  
(s1.id = s2.id+1 AND s1.id = s3.id+2)  
WHERE s1.people >= 100 AND s2.people >= 100 AND s3.people>=100  
ORDER BY visit_date
```



## 602. Friend Requests II: Who Has the Most Friends | Medium | [LeetCode](#)

In social network like Facebook or Twitter, people send friend requests and accept others' requests as well. Table `request_accepted` holds the data of friend acceptance, while `requesterid` and `accepterid` both are the id of a person.

TEXT

requester_id	accepter_id	accept_date
1	2	2010-03-03
1	3	2016-06-08

6 0

2010-03-03

2016-06-08

2	3	2016-06-08
3	4	2016-06-09

Write a query to find the the people who has most friends and the most friends number. For the sample data above, the result is:

TEXT

id	num	
----- -----		
3	3	

Note:

It is guaranteed there is only 1 people having the most friends. The friend request could only been accepted once, which mean there is no multiple records with the same requesterid and accepterid value. Explanation: The person with id '3' is a friend of people '1', '2' and '4', so he has 3 friends in total, which is the most number than any others.

Follow-up: In the real world, multiple people could have the same most number of friends, can you find all these people in this case?

SQL

```
SELECT t.id, sum(t.num) AS num
FROM (
    (SELECT requester_id AS id, COUNT(1) AS num
     FROM request_accepted
     GROUP BY requester_id)
    union all
    (SELECT accepter_id AS id, COUNT(1) AS num
     FROM request_accepted
     GROUP BY accepter_id)) AS t
GROUP BY t.id
ORDER BY num DESC
LIMIT 1;
```



## 603. Consecutive Available Seats | Easy | [LeetCode](#)

Several friends at a cinema ticket office would like to reserve consecutive available seats. Can you help to query all the consecutive available seats order by the seat\_id using the following cinema table?

TEXT

seat_id	free
1	1
2	0
3	1

6 | 0

	4		1	
	5		1	

Your query should return the following result for the sample case above.

TEXT

	seat_id	
-----		

6 | 0

	3	
	4	
	5	

Note:

The seat\_id is an auto increment int, and free is bool ('1' means free, and '0' means occupied.). Consecutive available seats are more than 2(inclusive) seats consecutively available.

## Solution

SQL

```
SELECT DISTINCT t1.seat_id
FROM cinema AS t1 JOIN cinema AS t2
ON abs(t1.seat_id-t2.seat_id)=1
WHERE t1.free='1' AND t2.free='1'
ORDER BY t1.seat_id
```



## 607.Sales Person | Easy | [LeetCode](#)

### Description

Given three tables: `salesperson`, `company`, `orders`. Output all the names in the table `salesperson`, who didn't have sales to company 'RED'.

### Example Input

Table: salesperson

TEXT

sales_id	name	salary	commission_rate	hire_date
1	John	100000	6	4/1/2006
2	Amy	120000	5	5/1/2010
3	Mark	65000	12	12/25/2008
4	Pam	25000	25	1/1/2005
5	Alex	50000	10	2/3/2007

The table salesperson holds the salesperson information. Every salesperson has a sales\_id and a name. Table: company

TEXT

com_id	name	city	6	0
1	RED	Boston		

	2	ORANGE	New York
	3	YELLOW	Boston
	4	GREEN	Austin
+-----+-----+-----+			

The table company holds the company information. Every company has a com\_id and a name. Table: orders

TEXT

order_id	date	com_id	sales_id	amount
1	1/1/2014	3	4	100000
2	2/1/2014	4	5	5000
3	3/1/2014	1	1	50000
4	4/1/2014	1	4	25000

The table orders holds the sales record information, salesperson and customer company are represented by *salesid* and *comid*. output

```
TEXT
+-----+
| name |
+-----+
| Amy   |
| Mark  |
| Alex  |
+-----+
```

## Explanation

According to order '3' and '4' in table orders, it is easy to tell only salesperson 'John' and 'Alex' have sales to company 'RED', so we need to output all the other names in table salesperson.

## Solution

```
SQL
SELECT name
FROM salesperson
WHERE name NOT IN
(SELECT DISTINCT salesperson.name
FROM salesperson, orders, company
WHERE company.name = 'RED'
AND salesperson.sales_id = orders.sales_id
AND orders.com_id = compa
  6 | 0
```

## 608. Tree Node | Medium | [LeetCode](#)

Given a table tree, id is identifier of the tree node and p\_id is its parent node's id.

TEXT

+-----+

| id | p\_id |

+-----+

 6 |  0

	1	null
	2	1
	3	1
	4	2
	5	2
+-----+		

Each node in the tree can be one of three types:

Leaf: if the node is a leaf node. Root: if the node is the root of the tree. Inner: If the node is neither a leaf node nor a root node. Write a query to print the node id and the type of the node. Sort your output by the node id. The result for the above sample is:

TEXT
+-----+
id   Type
+-----+
1   Root
2   Inner
3   Leaf
4   Leaf
5   Leaf
+-----+

## Explanation

Node '1' is root node, because its parent node is NULL and it has child node '2' and '3'. Node '2' is inner node, because it has parent node '1' and child node '4' and '5'. Node '3', '4' and '5' is Leaf node, because they have parent node and they don't have child node. And here is the image of the sample tree as below:

TEXT

```
1
 /   \
2     3
 /   \
4     5
```

## Note

If there is only one node on the tree, you only need to output its root attributes.

## Solution

6 | 0

SQL

44. Oracle Tree - SELECT TOTAL



```
## Basic ideas: LEFT JOIN
# In tree, each node can only one parent or no parent
## | id | p_id | id (child) |
## |-----+-----+-----|
## | 1 | null | ..... 1 |
## | 1 | null | ..... 2 |
## | 2 | .. 1 | ..... 4 |
## | 2 | .. 1 | ..... 5 |
## | 3 | .. 1 | ..... null |
## | 4 | .. 2 | ..... null |
## | 5 | .. 2 | ..... null |

SELECT t1.id,
CASE
    WHEN ISNULL(t1.p_id) THEN 'Root'
    WHEN ISNULL(MAX(t2.id)) THEN 'Leaf'
    ELSE 'Inner'
END AS Type
FROM tree AS t1 LEFT JOIN tree AS t2
ON t1.id = t2.p_id
GROUP BY t1.id, t1.p_id
```

## 610. Triangle Judgement | Easy | 🔒 LeetCode

A pupil Tim gets homework to identify whether three line segments could possibly form a triangle. However, this assignment is very heavy because there are hundreds of records to calculate. Could you help Tim by writing a query to judge whether these three sides can form a triangle, assuming table triangle holds the length of the three sides x, y and z.

TEXT

x	y	z
13	15	30
10	20	15

For the sample data above, your query should return the follow result:

TEXT

x	y	z	triangle
13	15	30	No
10	20	15	Yes

## Solution

SQL

6 | 0

SELECT x, y, z,



```
CASE
WHEN x+y>z AND y+z>x AND x+z>y THEN 'Yes'
ELSE 'No'
END AS triangle
FROM triangle
```

## 612. Shortest Distance in a Plane | Medium | [LeetCode](#)

Table point\_2d holds the coordinates (x,y) of some unique points (more than two) in a plane. Write a query to find the shortest distance between these points rounded to 2 decimals.

TEXT

x	y
-1	-1
0	0
-1	-2

The shortest distance is 1.00 from point (-1,-1) to (-1,2). So the output should be:

TEXT

shortest
1.00

Note: The longest distance among all the points are less than 10000.

## Solution

SQL

```
SELECT ROUND(MIN(SQRT((t1.x-t2.x)*(t1.x-t2.x) + (t1.y-t2.y)*(t1.y-t2.y))), 2) as shortest  
FROM point_2d AS t1, point_2d AS t2  
WHERE t1.x!=t2.x OR t1.y!=t2.y  
  
# SELECT ROUND(SQRT((t1.x-t2.x)*(t1.x-t2.x) + (t1.y-t2.y)*(t1.y-t2.y)), 2) AS shortest  
# FROM point_2d AS t1, point_2d AS t2  
# WHERE t1.x!=t2.x OR t1.y!=t2.y  
# ORDER BY shortest ASC  
# LIMIT 1
```

## 613. Shortest Distance in a Line | Easy | [LeetCode](#)

 6 |  0

Table point holds the x coordinate of some points on x-axis in a plane, which are all integers. Write a query to find the shortest distance between two points in these points.

TEXT

	x	
-----		
	-1	
	0	
	2	

6 | 0

The shortest distance is '1' obviously, which is from point '-1' to '0'. So the output is as below:

```
TEXT
| shortest|
|-----|
| 1      |
```

Note: Every point is unique, which means there is no duplicates in table point.

Follow-up: What if all these points have an id and are arranged from the left most to the right most of x axis?

## Solution

```
SQL
SELECT t1.x-t2.x AS shortest
FROM point AS t1 JOIN point AS t2
WHERE t1.x>t2.x
ORDER BY (t1.x-t2.x) ASC
LIMIT 1
```

**614. Second Degree Follower | Medium | 🔒 [LeetCode](#)**

In facebook, there is a follow table with two columns: followee, follower.

Please write a sql query to get the amount of each follower's follower if he/she has one.

For example:

TEXT	
+-----+-----+	
followee	follower
+-----+-----+	
A	B
B	C
B	D
D	E
+-----+-----+	

should output:

TEXT	
+-----+-----+	
follower	num
+-----+-----+	
	 6    0
+-----+-----+	

	B	2	
	D	1	

Explanation: Both B and D exist in the follower list, when as a followee, B's follower is C and D, and D's follower is E. A does not exist in follower list.

Note: Followee would not follow himself/herself in all cases. Please display the result in follower's alphabet order.

## Solution



SQL

```
## Explain the business logic
## . A follows B. Then A is follower, B is followee
## What are second degree followers?
## . A follows B, and B follows C.
## . Then A is the second degree followers of C
```

```
SELECT f1.follower, COUNT(DISTINCT f2.follower) AS num
FROM follow AS f1 JOIN follow AS f2
ON f1.follower = f2.followee
GROUP BY f1.follower;
```

## 615. Average Salary: Departments VS Company | Hard |

[LeetCode](#)

Given two tables as below, write a query to display the comparison result (higher/lower/same) of the average salary of employees in a department to the company's average salary. Table: salary

TEXT

id   employee_id   amount   pay_date
----- ----- ----- -----
1   1   9000   2017-03-31
2   2   6000   2017-03-31
3   3   10000   2017-03-31
4   1   7000   2017-02-28
5   2   6000   2017-02-28
6   3   8000   2017-02-28

The *employeeid* column refers to the *employeeid* in the following table *employee*.

TEXT

 6 |  0

employee_id   department_id
----- -----

1	1	
2	2	
3	2	

So for the sample data above, the result is:

6 | 0

TEXT

| pay\_month | department\_id | comparison |

2017-03	1	higher	
2017-03	2	lower	
2017-02	1	same	
2017-02	2	same	

Explanation In March, the company's average salary is  $(9000+6000+10000)/3 = 8333.33$ ... The average salary for department '1' is 9000, which is the salary of employeeid '1' since there is only one employee in this department. So the comparison result is 'higher' since  $9000 > 8333.33$  obviously. The average salary of department '2' is  $(6000 + 10000)/2 = 8000$ , which is the average of employeeid '2' and '3'. So the comparison result is 'lower' since  $8000 < 8333.33$ . With the same formula for the average salary comparison in February, the result is 'same' since both the department '1' and '2' have the same average salary with the company, which is 7000.

## Solution

SQL

```

SELECT t1.pay_month, t1.department_id,
       (CASE WHEN t1.amount = t2.amount THEN 'same'
             WHEN t1.amount > t2.amount THEN 'higher'
             WHEN t1.amount < t2.amount THEN 'lower' END) AS comparison
FROM
       (SELECT left(pay_date, 7) AS pay_month, department_id, avg(amount) AS amount
        FROM salary JOIN employee
        ON salary.employee_id = employee.employee_id
        GROUP BY pay_month, department_id
        ORDER BY pay_month DESC, department_id) AS t1
       JOIN
       (SELECT left(pay_date, 7) AS pay_month, avg(amount) AS amount
        FROM salary JOIN employee

```

```
ON salary.employee_id = employee.employee_id  
GROUP BY pay_month) AS t2  
ON t1.pay_month = t2.pay_month
```

□

## 618. Students Report By Geography | Hard | 🔒 [LeetCode](#)

A U.S graduate school has students from Asia, Europe and America. The students' location information are stored in table student as below.

TEXT

name	continent
Jack	America
Pascal	Europe
Xi	Asia
Jane	America

6 | 0

Pivot the continent column in this table so that each name is sorted alphabetically and displayed underneath its corresponding continent. The output headers should be America, Asia and Europe respectively. It is guaranteed that the student number from America is no less than either Asia or Europe. For the sample input, the output is:

TEXT

America	Asia	Europe
Jack	Xi	Pascal
Jane		

Follow-up: If it is unknown which continent has the most students, can you write a query to generate the student report?

## Solution

ANSWER

SQL

```
SELECT t1.name AS America, t2.name AS Asia, t3.name AS Europe
FROM
    (SELECT (@cnt1 := @cnt1 + 1) AS id, name
     FROM student
     CROSS JOIN (SELECT @cnt1 := 0) AS dummy
     WHERE continent='America'
     ORDER BY name) AS t1
LEFT JOIN
    (SELECT (@cnt2 := @cnt2 + 1) AS id, name
     FROM student
     CROSS JOIN (SELECT @cnt2 := 0) AS dummy
     WHERE continent='Asia'
     ORDER BY name) AS t2
ON t1.id = t2.id
LEFT JOIN
    (SELECT (@cnt3 := @cnt3 + 1) AS id, name
     FROM student
     CROSS JOIN (SELECT @cnt3 := 0) AS dummy
     WHERE continent='Europe'
     ORDER BY name) AS t3
ON t1.id = t3.id
```

## 619. Biggest Single Number | Easy | 🔒 [LeetCode](#)

Table number contains many numbers in column num including duplicated ones. Can you write a SQL query to find the biggest number, which only appears once.

14 6 | 14 0

TEXT

+---+

| num |  
+---+  
| 8 |  
| 8 |  
| 3 |  
| 3 |  
| 1 |  
| 4 |

6 | 0

| 5 |

For the sample data above, your query should return the following result:

```
TEXT
+---+
| num |
+---+
| 6 |
```

Note: If there is no such number, just output null.

## Solution

```
SQL 
SELECT IFNULL((
    SELECT num
    FROM number
    GROUP BY num
    HAVING count(1) = 1
    ORDER BY num DESC
    LIMIT 0, 1), NULL) AS num
```

## 620. Not Boring Movies | Easy | [LeetCode](#)

X city opened a new cinema, many people would like to go to this cinema. The cinema also gives out a poster indicating the movies' ratings and descriptions. Please write a SQL query to output movies with an odd numbered ID and a description that is not 'boring'. Order the result by rating.

For example, table `cinema`:

TEXT

	id	movie	description	rating	
	1	War	great 3D	8.9	
	2	Science	fiction	8.5	
	3	irish	boring	6.2	
	4	Ice song	Fantacy	8.6	
	5	House card	Interesting	9.1	

6 | 0

For the example above, the output should be:

TEXT

id   movie   description   rating
5   House card   Interesting   9.1
1   War   great 3D   8.9

## Solution

SQL

```
SELECT *
FROM Cinema
WHERE description <> 'boring' AND ID % 2 = 1
ORDER BY rating DESC;
```



## 626. Exchange Seats | Medium | [LeetCode](#)

Mary is a teacher in a middle school and she has a table `seat` storing students' names and their corresponding seat ids.

The column `id` is continuous increment.

Mary wants to change seats for the adjacent students.

Can you write a SQL query to output the result for Mary?

TEXT

	id	student
	1	Abbot
	2	Doris
	3	Emerson
	4	Green
	5	Jeames

For the sample input, the output is:

TEXT

	id	student
	1	Doris
	2	Abbot
	3	Green
	4	Emerson
	5	Jeames

6 | 0

```
| 5 | James |  
+-----+-----+
```

**Note:**

If the number of students is odd, there is no need to change the last one's seat.

## Solution

SQL

```
SELECT  
IF(id < (SELECT MAX(id) FROM seat), IF(id%2=0, id-1, id+1), IF(id%2=0, id-1, id)) AS id
```

```
FROM seat  
ORDER BY id;
```

=====

## 627. Swap Salary | [LeetCode](#)

Table: salary

TEXT

Column Name	Type
id	int
name	varchar
sex	ENUM
salary	int

id is the primary key for this table.

The sex column is ENUM value of type ('m', 'f').

The table contains information about an employee.

Write an SQL query to swap all 'f' and 'm' values (i.e., change all 'f' values to 'm' and vice versa) with a single update statement and no intermediate temp table(s).

Note that you must write a single update statement, DO NOT write any select statement for this problem.

The query result format is in the following example:

TEXT

Salary table:

id	name	sex	salary
1	A	m	2500

60

	2		B		f		1500	
	3		C		m		5500	
	4		D		f		500	
+-----+-----+-----+-----+								

Result table:

+-----+-----+-----+-----+

 6 |  0

| id | name | sex | salary |

```

+-----+-----+-----+
| 1   | A    | f    | 2500  |
| 2   | B    | m    | 1500  |
| 3   | C    | f    | 5500  |
| 4   | D    | m    | 500   |
+-----+-----+-----+
(1, A) and (2, C) were changed from 'm' to 'f'.
(2, B) and (4, D) were changed from 'f' to 'm'.

```

## Solution

SQL



# With IF

```
UPDATE Salary SET sex = IF(sex='m', 'f', 'm')
```

# With CASE

```
UPDATE Salary SET sex = CASE WHEN sex='m' THEN 'f' ELSE 'm' END
```

## 1045. Customers Who Bought All Products | Medium | 🔒

[LeetCode](#)

Table: Customer

TEXT

```

+-----+-----+
| Column Name | Type      |
+-----+-----+
| customer_id | int      |

```

```
| product_key | int      |  
+-----+-----+
```

product\_key is a foreign key to Product table. Table: Product

TEXT

```
+-----+-----+  
| Column Name | Type      |  
+-----+-----+  
| product_key | int      |  
+-----+-----+
```

product\_key is the primary key column for this table.

Write an SQL query for a report that provides the customer ids from the Customer table that bought all the products in the Product table.

For example:

 6 |  0

TEXT

Customer table:

customer_id	product_key
1	5
2	6
3	5
3	6
1	6

Product table:

product_key
5
6

Result table:

customer_id
1
3

The customers who bought all the products (5 and 6) are customers with id 1 and 3



## Solution

SQL

```
SELECT customer_id
FROM Customer
GROUP BY customer_id
HAVING count(DISTINCT product_key) = (
    SELECT count(1)
    FROM Product)
```



## 1050. Actors and Directors Who Cooperated At Least Three Times | Easy | [LeetCode](#)

Table: ActorDirector

TEXT

Column Name	Type
actor_id	int
director_id	int

60